Development verification record — 2026-10-17

These checks establish implementation behavior, not comparative sample quality. The primary CIFAR-10 study was incomplete when this record was written.

Sequence and source provenance

The public repository and pre-specified protocol were pushed in commit `55dc2e4`. Core numerical/training implementation (`2c9a45d`) and CLI (`a3dd622`) were pushed before the first MNIST optimizer run. The first GPU run used `a3dd622`. Subsequent evaluation/deadline support was pushed in `5b69c90`; the fixed-batch/pilot scheduler was pushed in `b0eb1f5`. Full reporting/export support was pushed in `af53767`. Git history and run `source_revisions` retain the exact full revisions.

The originating user instructions and implementation plan specified local training, a 24 GPU-hour ceiling, public code/checkpoints, and exploratory paired inference. No runtime LLM was introduced. The implementation used TorchCFM, torchdiffeq, torch-fidelity, PyTorch, POT, and SciPy instead of recreating their maintained algorithms. AKOrN source had no repository license during inspection, so its published equations were implemented independently.

Experiments and observations

Check	Observed result	Interpretation
MNIST GPU smoke, 50 updates	Completed; about 17.68 training seconds; peak allocated memory 203,779,584 bytes	CUDA forward/backward/optimizer/checkpoint path executes
Resume to 100 updates	Completed; cumulative training time 33.96 seconds; EMA validation velocity MSE 2.00869	Saved run state can continue on GPU; early EMA remains close to initialization
Exact CPU resume test	Resumed and uninterrupted model/EMA/noise-generator states match exactly	Optimizer-boundary persistence is complete in the tested CPU setting
Fixed batch, 500 updates	MSE 2.255753 → 0.000602648; ratio 0.000267161	Full oscillator model can learn fixed production-path velocity targets
Same-input determinism	Exact repeated predictions in the fixed-batch GPU check	No hidden state is carried across these forward calls
Geometry	Maximum norm error 5.96e-8; maximum tangent inner-product magnitude 3.05e-7	FP32 spherical operations meet the predeclared 1e-5 numerical tolerance
CPU sampling	16 samples at 32 NFE completed; full float outputs and grid saved	GPU is not required for inference
GPU sampling/evaluation smoke	128 samples at 32 NFE completed; MNIST pixel SWD 1.101009	Full evaluation path works; this noise-like early model is not a quality result
Portable checkpoint reload	Original and safetensors-export CPU samples exactly equal; maximum difference zero	Safe export preserves the actual inference function
Budget enforcement	Isolated deliberately blocked worker was terminated; reservation retained	The ceiling is enforced independently of cooperative model code
Report/PDF workflow	Markdown and PDF records generated; first page and equations inspected	Pandoc/Typst artifacts exist and render the mathematical method
Planning-power reproduction	0.2406 at standardized effect 0.8 using 20,000 Gaussian simulations	Reproduces the pre-training low-power warning; not observed empirical power

The CPU preview after 100 optimizer updates visually remained noise-like. This was expected from a very short development run and a slowly updated EMA; it was retained rather than selected away. A longer fixed-time pilot, not an extension based on significance, was already part of the protocol.

Issues, decisions, and limits

The initial resume invocation omitted the required path after `--resume`; argparse rejected it before computation. The corrected documented form is `ofm train --resume <run-directory>`. Import-order lint findings were repaired while implementation progressed; CI run 38977501341 failed only an import-format check, repaired in `549327f`. A third-party POT SyntaxWarning did not affect numerical tests. No nonfinite-loss or gradient failure was observed in the checks above. A CLI help traceback was removed, and the scheduler was hardened to load actual frozen pilot settings instead of allowing later environment-default drift.

Tests for missing flow/training/evaluation/reporting/publication interfaces were introduced before those implementations, and their initial collection failures were observed. Some initial model tests could only execute after CUDA dependencies finished installing; this record does not claim every line followed a strict observed red-green sequence. The initial complete local suite contained 27 passing tests; additional help/freeze regression tests were added for the two later repairs. CI must pass on the merged implementation.

The 20-minute-per-architecture MNIST pilot uses the pre-specified optimizer settings and validates generalization/visual behavior before freezing CIFAR settings. No CIFAR result was used to choose these parameters. A fixed per-update EMA decay and differing architecture throughput are potential comparative confounders; they are reported as limitations rather than silently adjusted after outcomes. GPU utilization and system background load are not hardware-independent compute measures.

Sources: protocol, implementation contracts, Karpathy's training recipe, PyTorch serialization, SciPy permutation tests.