Failure: checks/inception-cache-failure/failure

Purpose and interpretation

This attempt did not complete normally. Its existence is retained; missing metrics are not replaced by zero or silently removed from the planned study. Diagnose the recorded error before any authorized retry.

Recorded evidence


```
{
  "status": "failed",
  "kind": "CPU evaluator-cache preparation",
  "error": "ModuleNotFoundError: No module named 'ofm.terminal'",
  "cause": "A read/import command overlapped the transient checkout performed
by gh pr merge --delete-branch.",
  "resolution": "Waited for the merge and checkout to finish, then reran the
unchanged cache preparation successfully. Future merges during active runs
retain the current checkout.",
  "gpu_seconds": 0,
  "id": "checks/inception-cache-failure/failure",
  "artifact": "runs/checks/inception-cache-failure/failure.json"
}
```


Method choices and learning

The shared optimizer, data policy, architecture, and stopping rule were chosen before final comparison and are described in the technical paper. The 3e-4 learning rate is a conservative AdamW starting point, not a tuned optimum. Microbatches bound laptop memory; four internal updates bound recurrent cost; EMA smooths noisy parameter updates. These choices trade study breadth for a fixed local compute budget. This record supports only the measurements above; causes of quality differences require the trained controls and uncertainty analysis. Full trajectories, populations, RNG states, and logs remain in the referenced artifact bundle. Sources: protocol, TorchCFM, AKOrN.